



Jenkins and Chef

Infrastructure CI and Application Deployment

Dan Stine
Copyright Clearance Center
www.copyright.com

June 18, 2014

#jenkinsconf

About Me

- Software Architect
- Library & Framework Developer
- Infrastructure Lead & Product Owner
- Enemy of Inefficiency and Needless Inconsistency



dstine at copyright.com

sw at stinemail.com

github.com/dstine

About Copyright Clearance Center

- Global licensing solutions that make © work for everyone
 - Get, share, and manage content
 - Rights broker for world's most sought after materials
 - Global company (US, Europe, Asia) – HQ in Danvers, MA
- Industry-specific software systems
 - Internal and external user base
 - Applications, services, databases
 - Organic growth over many years
- In 2011, CCC adopted a Product Platform strategy for growing its software portfolio



Agenda

- Context
- Primer
- Deployment Process Design
- Cookbook Builds
- Application Deployment
- Wrap Up






CONTEXT

Standard Software Platform

- Started platform definition in 2011
 - Homogeneous by default
- Tools
 - Java, Spring, Tomcat, Postgres
 - Git/GitHub, Gradle, Jenkins, Artifactory, Liquibase
- Process
 - Standard development workflow
 - Standard application shape & operational profile



Initial Delivery Pipeline



Initial Delivery Pipeline

- Automated build process
- Publish build artifacts to Artifactory
 - Application WARs
 - Liquibase JARs
- Manual deploys
 - (Many apps) x (many versions) x (multiple environments) = TIME & EFFORT
 - The more frequently a task is performed, the greater the return from improved efficiency





Improved Deployment Process

- Goals
 - Reduce effort
 - Improve speed, reliability, and frequency
 - Handle app deploys and db schema updates
 - Enable self-service
- Process Changes
 - Manual → Automated
 - Prose instructions → Infrastructure as code



Target Delivery Pipeline






PRIMER

Layers of System Management

- Orchestration
 - Processes collaborating in a distributed system
- Configuration
 - Install and configure packages and software
- Provisioning
 - Hypervisor (VMware, EC2)



Infrastructure as Code

- Develop and manage software infrastructure with practices similar to those used to develop software applications
- Examples
 - Source Code
 - Modularity
 - Abstraction
 - Testing



Configuration Management


“Process for establishing and maintaining consistency of a product’s performance, functional and physical attributes with its requirements, design and operational information throughout its life” ([wikipedia](#))



- CM tools for managing software systems
 - CFEngine, Puppet, Chef, Salt, Ansible
- Embody Infrastructure as Code principles
- Define desired state of machine
 - Each run inspects state and makes necessary changes, if any

Chef

- Configuration management tool
- Exposes DSL hosted in Ruby
 - Express “what” not “how”
 - Clean, purposeful capture of intent
- Favor DSL when writing code
 - Ruby is available, if required



Chef Terminology (1)

- Chef client is installed on nodes (machines) which are registered with the Chef server
- Developers write code on workstations and use tools such as knife to interact with server
- Chef models node configuration as a set of DSL resources (e.g. package, service, directory) which are mapped to internal providers (actual code to execute)
 - Can define custom resources



Example Chef Code

This resource declaration

```
directory '/a/b/c' do
  owner 'admin'
  group 'admin'
  mode '0755'
  action :create
  recursive true
end
```


ensures that

```
$ ls -ld /a/b/c
drwxr-xr-x. 5 admin admin 4096 Feb 14 11:22 /a/b/c
```



Chef Terminology (2)

- A recipe declares a set of resources with desired configuration
- A cookbook contains a set of semantically-related code and is the fundamental unit of distribution for Chef code
 - Compare to JAR for Java code
- A data bag holds JSON information in one or more data bag items accessible from Chef code
- Chef environments model deployed environments
- Each node has a run list containing recipes





DEPLOYMENT PROCESS DESIGN

Basic Approach

- Deploy custom applications with Chef
- Execute schema updates with Liquibase
- Coordinate everything with:



Jenkins



Jenkins as Coordinator

- General purpose job executor
 - Shell script steps for Chef API
 - Gradle steps for Liquibase updates
 - Arbitrary code at any point in lifecycle
- UI
 - Smooth integration with Active Directory
 - Authentication and authorization
- Administration
 - Familiar with Jenkins from application build jobs



CCC Application Group

- Set of deployable units that are versioned and released together
- For example, a group might have
 - UI
 - REST service
 - Message consumer
 - DB
- Build with a single command
- Deploy with a single command



Technical Design Goals

- Provide clean API
 - Specify only essential differences between apps
 - Custom Chef resource is the interface
 - Codify & enforce standards
- Balance consistency & flexibility
 - Code in semantically-versioned cookbooks
 - Configuration in data bags
- Controlled cookbook promotion
 - Chef environment specifies cookbook version constraint



Cookbook Types

- Library Cookbooks
 - Encapsulate common re-usable logic
 - Define custom resource to install an app
 - And the implementing provider
- Application Cookbooks
 - Depend on library cookbooks
 - One cookbook per application group
 - One recipe per application
 - Recipes use custom resource
 - Lightweight



Data Bags

- Contain application configuration
 - Service endpoints, JAVA_OPTS, etc.
- One data bag per application group
 - One data bag item per environment
- “Live” reflection of deployed configuration
 - Edit → push to Chef server → deploy
 - Master always matches state of Chef server



Custom Resource Usage


Ensure my-app-ui WAR is deployed:

```
ccc_webapp "my-app-ui" do
  provider :ccc_webapp

  artifact_group           'com.copyright.myapp'
  artifact                 'my-app-ui'

  container                'MY-APP-UI'
  http_port                '8080'
  shutdown_port             '8900'

  properties_template       'ccc.properties.erb'
  app_group_data_bag        'my_app'
end
```



Custom Resource Actions

- Retrieves Java, Tomcat & WAR from Artifactory
- Installs Java and Tomcat in standard locations
- Creates and configures Tomcat container
- Installs WAR in the container
- Opens port in host firewall
- Generates application properties file
- Starts container

(Each action taken only if necessary)



Data Bag Structure

data_bags/my_app/DEV.json (data bag item)

```
"version": "1.4.9",
"runtime": {
    "my-app-ui": {
        "java_opts": "-Xmx2G -XX:MaxPermSize=1024m"
    }
},
"app_config": {
    "db.url":      "jdbc:postgresql://devdb:5432/myapp",
    "svc.foo.url": "http://devsvc:9000/foo"
}
```

data_bags/my_app/TEST.json

...

data_bags/my_app/PROD.json



Cookbook Data Bag Code *

ccc/providers/webapp.rb (library cookbook)

```
app_group_data = data_bag_item(app_group_data_bag,  
node.chef_environment)
```

```
java_opts = app_group_data['runtime'][artifact]['java_opts']  
// pass java_opts to Tomcat container
```

```
app_config = app_group_data['app_config']  
// pass app_config to template resource declaration
```

my_app/templates/ccc.properties.erb (application cookbook)

```
db.url=<%= @app_config['db.url'] %>  
svc.foo.url=<%= @app_config['svc.foo.url'] %>
```




* Included for future reference

Roles


- Deployers
 - Update data bags & environment files
 - Initiate deployments
- Tech leads
 - Maintain application cookbooks
- Framework developers
 - Maintain library cookbooks
 - Maintain framework
 - Process improvement





COOKBOOK BUILDS

Cookbook Build Process




Jenkins Build Server

- For each application group
 - Cookbook CI job
 - Cookbook release job
 - Same master as application build jobs
- New class of slaves
 - Ruby with required gems
 - Chef with credentials for Chef server
 - EC2 credentials to create test nodes




Cookbook CI Job

- Triggered when new Chef code is merged
- Static analysis
 - JSON syntax (json gem)
 - Ruby syntax and style (Tailor)
 - Chef syntax (Knife)
 - Chef style and correctness (Foodcritic)
- Integration testing
 - Test Kitchen with kitchen-ec2 plugin




Integration Testing Lifecycle

- Spin up EC2 instance(s) to mimic actual deployment topology of application group
- Run Chef on each instance (node)
- Execute asserts – pass or fail
- Dispose of instance(s)



Integration Testing Details

- Instances created from AMI
 - Preconfigured with Ruby and Chef
- Using Chef Solo
 - Avoid adding ephemeral nodes to Chef server
- Faux Chef environment “CHEFDEV”
 - JSON for real environments is reserved
- Tag EC2 instances for traceability
- Troubleshoot by running Test Kitchen from workstation



Cookbook Release Job


- Triggered manually
- Runs same tests as CI job
- Uploads new cookbook version to Chef server
- Tags Git repo





APPLICATION DEPLOYMENT

Application Deploy Process



Jenkins Deploy Server

- Separate master for deploys
- Slaves
 - Ruby with required gems
 - Chef with credentials for Chef server
 - SSH keys for nodes



Deploy Job Types

- Each app group has two deploy jobs
 - DEV deploy for Development
 - Non-DEV deploy for Operations
 - Will provide more flavors over time
- Job parameters
 - Environment (non-DEV jobs only)
 - Application group version



What Does a Deployer Do?

- Makes configuration changes
 - Edits application data bag item
 - Edits environment file (if necessary)
 - Merges code
- Executes job in Jenkins




Example Deploy Job Run

- Deployer enters parameters
 - Application version = 1.4.9
 - Environment = TEST
- Then automation takes over
 - Confirms my_app data bag has TEST version 1.4.9
 - Uploads TEST environment file and my_app data bag item for TEST to Chef server
 - Finds all nodes in TEST environment with run list containing my_app recipes
 - Runs Chef client on each found node
 - Sends email notification



Push-Button Deploys



Project rs.DEV1

This build requires parameters:

APP_GROUP_VERSION

Application group version to deploy. Must match databag value.

Project rs

This build requires parameters:

ENVIRONMENT

Environment to which the application group should be deployed.

APP_GROUP_VERSION

Application group version to deploy. Must match databag value.

Deploy History

	Build History	(trend)
●	#33 May 29, 2014 3:37:06 PM [DEV1] 3.1.5	184 KB
●	#32 May 21, 2014 11:20:21 AM [DEV1] 3.1.4	245 KB
●	#31 May 14, 2014 4:02:46 PM [DEV1] 3.1.3	155 KB
●	#30 May 13, 2014 2:32:31 PM [DEV1] 3.1.2	149 KB
●	#29 May 13, 2014 12:00:44 PM [DEV1] 3.1.2	197 KB
●	#28 May 6, 2014 4:25:20 PM [DEV1] 2.1.29	199 KB
●	#27 May 1, 2014 4:04:00 PM [DEV1] 2.1.27	244 KB
●	#26 Apr 10, 2014 9:15:28 AM [DEV1] 2.1.27	132 KB
●	#25 Apr 9, 2014 11:12:55 AM [DEV1] 3.1.1	130 KB
●	#24 Apr 9, 2014 10:12:14 AM [DEV1] 3.1.0	129 KB
●	#23 Apr 9, 2014 10:02:18 AM [DEV1] 3.1.0	70 KB
●	#22 Apr 4, 2014 1:21:34 PM [DEV1] 2.1.26	157 KB
●	#50 May 30, 2014 5:47:41 PM [PRDC1] 2.1.31	128 KB
●	#49 May 30, 2014 5:38:36 PM [PRDC2] 2.1.31	129 KB
●	#48 May 30, 2014 3:42:07 PM [PS1] 2.1.31	137 KB
●	#47 May 30, 2014 12:47:01 PM [QA1] 2.1.31	135 KB
●	#46 May 29, 2014 10:28:04 AM [PRDC2] 2.1.30	124 KB
●	#45 May 29, 2014 10:06:02 AM [PRDC2] 2.1.30	145 KB
●	#44 May 29, 2014 9:53:16 AM [PRDC1] 2.1.30	166 KB
●	#43 May 22, 2014 11:05:57 AM [PS1] 2.1.30	128 KB
●	#42 May 22, 2014 10:56:09 AM [PS1] 2.1.30	289 KB
●	#41 May 15, 2014 2:00:35 PM [QA1] 2.1.30	222 KB
●	#40 May 9, 2014 11:32:42 AM [PREC1] 2.1.29	145 KB
●	#39 May 5, 2014 3:43:50 PM [QA1] 2.1.29	130 KB





WRAP UP

Most Important Advice

- Beware of overly prescriptive “lessons learned” and “best practices”
- Synthesize a solution for your context
- That said...




Principles & Guidelines (1)

- Standardize wherever possible
 - Technology, design, process
 - Achieve economies of scale
 - Exceptions are permissible but rare
- Every tool must have an API
 - Avoid “hitting the wall” down the road
 - Tradeoff some out-of-the-box capabilities




Principles & Guidelines (2)

- Use multiple communication paths
 - All-hands presentations
 - Kickoff meetings with each team
 - Developer walkthroughs
 - Documentation
- Be opportunistic
 - Find and nurture your early adopters



Principles & Guidelines (3)

- Balance process and progress
 - Must provide tangible results
 - And also build foundation for future
 - Just like with application development!
- Start with a big pain point
 - Providing relief builds credibility going forward
 - Hopefully recoups bandwidth to reinvest



“When Will You Be Done?”

- DONE is a dangerous word
 - Business won't stop evolving
 - Neither will the supporting applications
 - Nor should the supporting infrastructure
- X is a journey, not a destination
 - For many values of X
 - Deployment automation
 - Continuous delivery
 - DevOps



Thank You

- Copyright Clearance Center
Engineering Team
- Jenkins User Conference
Organizers and Sponsors



Resources

- “Infrastructure in the Cloud Era”

Adam Jacob and Ezra Zygmuntowicz

<http://www.slideshare.net/adamhjk/infrastructure-in-the-cloud-era>

<http://www.youtube.com/watch?v=HaABapTwQ2c>

- Chef cookbook versioning policy

<http://chef-community.github.io/cvp/>



Thank You To Our Sponsors

Platinum



Gold



MidVision™
Release the innovation



Silver

